

Adaptive Learning in Motion: Harnessing Cloud-Based AI for a Smart Soccer Ball's Real-Time Navigation

City University of Seattle
School of Technology & Computing
TEAM A

Verónica Elze
ElzeVeronica@CityUniversity.edu
Master of Artificial Intelligence

Vidhyalakshmi Amarnath
AmarnathVidhyalaksh@CityUniversity.edu
Master of Artificial Intelligence

Honorine Ndom Ndzah
NdomndzahHonorine@CityUniversity.edu
Master of Computer Science

Jayasudha Kalamegam
KalamegamJayasudha@CityUniversity.edu
Master of Computer Science

Abstract

This project aims to develop a Smart Soccer Ball, a compact, autonomous system that utilizes reinforcement learning, real-time sensor technology, and cloud-based AI platforms to navigate dynamic environments. Leveraging AWS SageMaker for RL model training and SimSpace Weaver for scalable 3D simulations, the soccer ball will be trained to adapt to both static and dynamic obstacles within a puzzle maze. The project will be conducted in two phases: Phase 1 focuses on developing a Minimum Viable Product with basic maze navigation and static obstacles, while Phase 2 (contingent on efficient progress) will explore advanced RL techniques, dynamic environments, and sophisticated sensors for complex obstacle handling. This endeavor provides an opportunity for hands-on AI and robotics experience, offering insights into low-cost, adaptable autonomous systems and smart sports technology. By demonstrating adaptive, real-time decision-making in a compact device, the project highlights the potential of applying advanced AI to sensor-driven robotics.

Keywords: smart soccer ball, reinforcement learning, AWS SageMaker, autonomous navigation, adaptive robotics, dynamic environments

1. INTRODUCTION

Problem to Solve

One of the key challenges we face is navigating dynamic environments with an autonomous system. Although autonomous navigation has been explored in various fields, it remains complex when environments constantly change. Reinforcement learning (RL) models, though powerful, often struggle to generalize across varying conditions and adjust to real-time changes (Kober et al., 2013). Traditional RL approaches, which rely heavily on trial and error, can be inefficient in unpredictable environments.

This is where model-based reinforcement learning (MBRL) steps in, enhancing decision-making by allowing systems to anticipate and adapt to changes. MBRL has proven effective in sensor-driven robotics, addressing real-time navigation challenges (Polydoros & Nalpantidis, 2017). By incorporating MBRL, we aim to improve the adaptability and efficiency of the Smart Soccer Ball, enabling it to tackle both static and dynamic obstacles.

Motivation

Teaching an autonomous agent to navigate dynamic environments presents exciting challenges with real-world implications. Our goal is to develop a Smart Soccer Ball capable of learning how to navigate a puzzle maze. RL models show promise but often struggle to adapt to ever-changing environments (Kober et al., 2013).

MBRL enhances decision-making and control by planning ahead, making it particularly useful in environments that require real-time adaptability (Polydoros & Nalpantidis, 2017). This method will allow our soccer ball to handle dynamic obstacles effectively.

There's also growing interest in low-cost autonomous systems in industries like smart sports technology, where AI-driven devices enhance performance (Zhang, 2022). By integrating sensors and leveraging AWS SageMaker for real-time decision-making, we aim to demonstrate how small-scale AI systems can efficiently tackle complex tasks (Malekzadeh et al., 2018). Plus, let's be honest, it's undeniably entertaining to see a soccer ball outwit a maze.

Usefulness/Beneficiaries

This project demonstrates practical applications for industries and education alike. By leveraging RL, we show how autonomous systems can adapt

to dynamic environments, a challenge for traditional models (Kober et al., 2013).

For industries like smart sports, MBRL offers valuable insights into developing low-cost systems that respond to environmental changes. It improves decision-making by anticipating rather than reacting (Polydoros & Nalpantidis, 2017).

Educationally, this project provides hands-on experience with sensor integration and AI platforms like AWS SageMaker, enabling real-time decision-making (Malekzadeh et al., 2018). These skills are essential as AI and robotics continue to grow.

Beyond sports, this project contributes to autonomous navigation research, showing how small-scale systems can handle complex tasks. Watching a soccer ball master a maze? Definitely a bonus (Zhang, 2022).

2. LITERATURE REVIEW

Autonomous systems capable of adapting to dynamic environments are increasingly relevant across various domains, including robotics, artificial intelligence (AI), and sports technology. The Smart Soccer Ball project aims to develop a compact, sensor-driven device that utilizes reinforcement learning (RL) to navigate mazes and dynamic obstacles. This literature review explores foundational research on similar projects, identifies applicable methodologies, and highlights how these insights contribute to the development of this project.

Reinforcement Learning in Robotics

Reinforcement learning has emerged as a transformative approach for robotics, enabling systems to learn optimal policies through trial and error. Kober et al. (2013) provide a comprehensive survey of RL applications in robotics, emphasizing its utility in environments requiring adaptive decision-making. Their research highlights challenges such as designing appropriate reward structures and balancing exploration versus exploitation—critical elements in the Smart Soccer Ball's ability to learn efficient navigation strategies. Unlike traditional rule-based systems, RL models dynamically adjust to new scenarios, offering flexibility that aligns with the project's goals.

The Smart Soccer Ball project specifically benefits from model-based reinforcement learning (MBRL). As Polydoros and Nalpantidis (2017)

explain, MBRL enhances a system's ability to predict and respond to environmental changes. By leveraging these predictive capabilities, the soccer ball can navigate dynamic mazes more effectively than model-free methods allow. This focus on adaptability differentiates the project from similar applications of RL in static environments.

Cloud Computing and AI Scalability

The integration of cloud-based platforms like AWS SageMaker is another pivotal aspect of the Smart Soccer Ball project. SageMaker provides a scalable environment for developing, training, and deploying RL models, offering computational resources that significantly reduce the time and cost of experimentation (Amazon Web Services, 2021). This capability contrasts with earlier robotics projects that relied heavily on on-premises hardware, which limited scalability.

Tripuraneni and Song (2019) underscore the role of cloud infrastructure in accelerating AI model development and deployment. Their work demonstrates how platforms like SageMaker streamline workflows by enabling distributed computing and iterative model tuning. This scalability is critical for the Smart Soccer Ball, where rapid experimentation with various RL algorithms—such as proximal policy optimization (PPO) and soft actor-critic (SAC)—is essential for optimizing navigation strategies (Schulman et al., 2017; Haarnoja et al., 2018).

Autonomous Navigation and Sensor Integration
Sensor-driven systems form the backbone of many autonomous robotics projects. Malekzadeh et al. (2018) show how integrating accelerometers and proximity sensors enhances navigation accuracy in compact robots. Their findings are directly applicable to the Smart Soccer Ball's design, where sensors will enable the ball to detect and avoid obstacles. Unlike projects that rely solely on vision-based systems, this project incorporates multiple sensor modalities to improve real-time decision-making.

The use of LIDAR as a potential upgrade aligns with research by Malla and Dholakiya (2022), who demonstrate its effectiveness in mapping dynamic environments. While LIDAR is not part of the initial phase of the Smart Soccer Ball project, its inclusion in future iterations could significantly enhance spatial awareness and obstacle detection capabilities. This step would elevate the system's performance to meet the challenges of complex, real-world scenarios.

Comparative Analysis with Similar Projects

Zhang (2022) explores AI applications in soccer training, employing genetic algorithms for optimizing movement paths. While genetic algorithms differ from RL, Zhang's work underscores the value of AI in dynamic navigation. The Smart Soccer Ball project builds on these insights by applying RL techniques that offer greater flexibility and adaptability. Unlike genetic algorithms, which evolve solutions over multiple iterations, RL models learn in real time, making them more suited for the dynamic obstacle scenarios envisioned for this project.

Another comparable effort is Shah et al.'s (2020) AirSim framework, which uses high-fidelity simulations for training autonomous vehicles. The Smart Soccer Ball leverages similar principles by incorporating AWS SimSpace Weaver for dynamic environment simulations. However, while AirSim focuses on vehicles with high computational demands, this project emphasizes compact, cost-effective robotics.

Integration and Contribution

The Smart Soccer Ball project synthesizes advances in reinforcement learning, cloud-based AI, and sensor-driven robotics to create an autonomous system with real-world applications. By leveraging model-based RL, scalable cloud platforms like AWS SageMaker, and advanced sensor technologies, the project addresses challenges in navigating dynamic environments. Compared to related work, the project stands out for its emphasis on cost-effective design and adaptability, demonstrating the potential of AI in sports and robotics.

This literature review highlights the foundational research and methodologies that guide the Smart Soccer Ball's development, positioning it as a pioneering effort in the intersection of autonomous systems and sports technology.

3. METHODOLOGY

To develop a reinforcement learning (RL) solution for navigating a simulated maze environment, we have focused on building and validating the foundational components locally. This approach ensures a controlled development environment before transitioning to cloud-based deployment and scaling.

Environment Setup and Customization

Using the MiniGrid library, we created a grid-based maze environment tailored for RL tasks. A

custom observation wrapper was implemented to preprocess and flatten the multi-dimensional observations into a single vector, streamlining compatibility with RL algorithms. Discrete actions, including movements in cardinal directions and interactions with the environment, were defined to enable efficient navigation by the agent. All customization and validation of the environment have been carried out in a local environment, ensuring that the setup is robust before integration with external systems.

Agent Training

The Proximal Policy Optimization (PPO) algorithm was used to train the agent, leveraging the Stable Baselines3 framework. Training was executed locally, utilizing the DummyVecEnv wrapper to vectorize the environment and enable efficient agent-environment interactions. This local setup allowed iterative testing and debugging in a controlled environment.

Training progress was monitored with TensorBoard, which provided real-time insights into key metrics such as rewards, policy loss, and training duration. By logging these metrics, we validated the agent's learning progress and fine-tuned parameters, such as learning rate and policy architecture, for improved performance. The agent's speed in reaching the green square goal showed a notable improvement, decreasing from an initial navigation time of approximately 106 seconds to just 2 seconds by the end of training.

Evaluation and Validation

Our local setup also supported detailed evaluation and debugging of the environment and agent behavior. Observations were validated to ensure they were processed correctly, and actions were tested to confirm expected outcomes. Debugging efforts focused on resolving compatibility issues, such as observation flattening errors and reward signal inconsistencies. The validation process demonstrated that the agent could learn and adapt within the customized environment.

This locally built foundation ensures that all components are well-tested and functional before deploying to AWS for further scalability and performance enhancements. By prioritizing local implementation, we have established a reliable framework to transition into more complex, cloud-based simulations.

Tools

Our Smart Soccer Ball project will leverage a combination of hardware, software, and cloud-based tools to facilitate reinforcement learning

(RL) and real-world agent interactions. These tools are integral to developing, testing, and deploying our adaptive navigation system.

Hardware (to be implemented)

- Sphero Mini App-Enabled Robotic Ball: Will serve as the physical embodiment of our RL agent, translating digital decision-making into tangible movements.
- Mobile Device: Will be used to control and monitor the Sphero Mini through the Sphero Edu App, ensuring seamless interaction during testing phases.

Software Tools

- Python Programming Language: Provides the foundation for RL model development, cloud integration, and environment customization.
- MiniGrid Library: A minimalistic, grid-based environment used for creating and customizing maze scenarios for training our RL agent.
- Stable Baselines3 Framework: Facilitates the implementation of the Proximal Policy Optimization (PPO) algorithm, chosen for its efficiency and reliability in RL tasks.
- OpenAI Gym: Allows preliminary testing in a simulated environment to refine RL strategies before deployment.

Cloud Tools (to be implemented)

- AWS SageMaker: Will handle the training and fine-tuning of RL models, leveraging its scalable infrastructure to accelerate development cycles.
- AWS SimSpace Weaver: Will simulate dynamic, three-dimensional puzzle mazes for testing agent adaptability and optimizing navigation strategies.

Development and Monitoring Tools

- TensorBoard: Enables real-time visualization of training metrics such as rewards, loss, and policy performance.
- GitHub: Will ensure version control and collaborative tracking of code and model iterations.

By integrating these tools, the project will achieve a balance between virtual simulation and real-world testing, fostering a robust development environment for adaptive RL models.

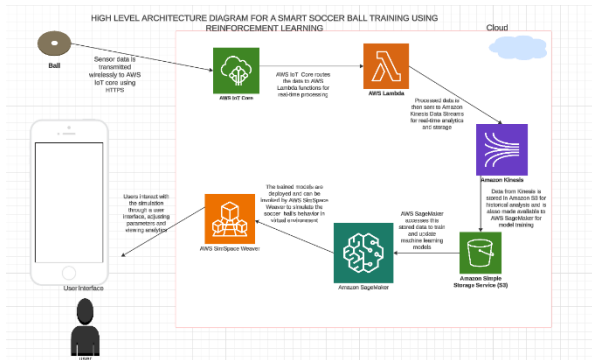


Figure 1 Architecture Diagram

High-Level Architecture Overview

The high-level architecture diagram (Figure 2, see also Appendix A) illustrates the proposed integration of hardware, cloud-based AI tools, and simulation environments in the Smart Soccer Ball project. This architecture enables seamless data flow, real-time analytics, and reinforcement learning model training. Key components include:

- **AWS IoT Core:** Facilitates wireless data transmission from the soccer ball's sensors to the cloud.
- **AWS Lambda:** Processes sensor data in real-time, triggering downstream workflows.
- **Amazon Kinesis and S3:** Streams processed data for real-time and historical analytics, stored for subsequent use.
- **AWS SageMaker:** Trains and refines reinforcement learning models using collected data.
- **AWS SimSpace Weaver:** Simulates dynamic environments, allowing the soccer ball to learn and adapt in a virtual 3D maze.

This architecture will leverage cloud computing for scalability, ensuring efficient training and evaluation of the soccer ball's navigation capabilities.

4. ACKNOWLEDGMENTS

I would like to acknowledge the assistance of OpenAI's ChatGPT tool in improving the verbiage, flow, and transitions of this paper.

5. REFERENCES

The following references include a combination of sources recommended by ChatGPT, with the majority published or updated within the last five years. All citations adhere to APA (American Psychological Association) guidelines.

Scholarly Sources

These sources provide foundational research and insights relevant to the project.

- Amazon Web Services. (2021). *AWS SageMaker: Developer Guide*. Retrieved from <https://docs.aws.amazon.com/sagemaker/latest/dg/whatis.html>
- Haarnoja, T., Zhou, A., Abbeel, P., & Levine, S. (2018). Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. *Proceedings of the 35th International Conference on Machine Learning (ICML)*, 2961–2970. <https://arxiv.org/abs/1801.01290>
- Kober, J., Bagnell, J. A., & Peters, J. (2013). Reinforcement learning in robotics: A survey. *The International Journal of Robotics Research*, 32(11), 1238–1274. <https://doi.org/10.1177/0278364913495721>
- Malla, A., & Dholakiya, M. (2022). Autonomous navigation in unstructured environments using LIDAR-based reinforcement learning. *International Journal of Advanced Robotic Systems*, 19(1), 1–12. <https://doi.org/10.1177/17298814221075422>
- Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., ... & Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529–533. <https://doi.org/10.1038/nature14236>
- Polydoros, A. S., & Nalpantidis, L. (2017). Survey of model-based reinforcement learning: Applications on robotics. *Journal of Intelligent & Robotic Systems*, 86(2), 153–173. <https://doi.org/10.1007/s10846-017-0570-0>
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. *arXiv Preprint*. <https://arxiv.org/abs/1707.06347>
- Shah, S., Dey, D., Lovett, C., & Kapoor, A. (2020). AirSim: High-fidelity visual and physical simulation for autonomous vehicles. *Conference on Autonomous Agents and Multiagent Systems (AAMAS)*. Retrieved from <https://github.com/microsoft/AirSim>
- Tripuraneni, V., & Song, M. (2019). Cloud infrastructure for scalable machine learning: A comprehensive study. *Journal of Cloud Computing*, 8(1), 1–18. <https://doi.org/10.1186/s13677-019-0124-9>

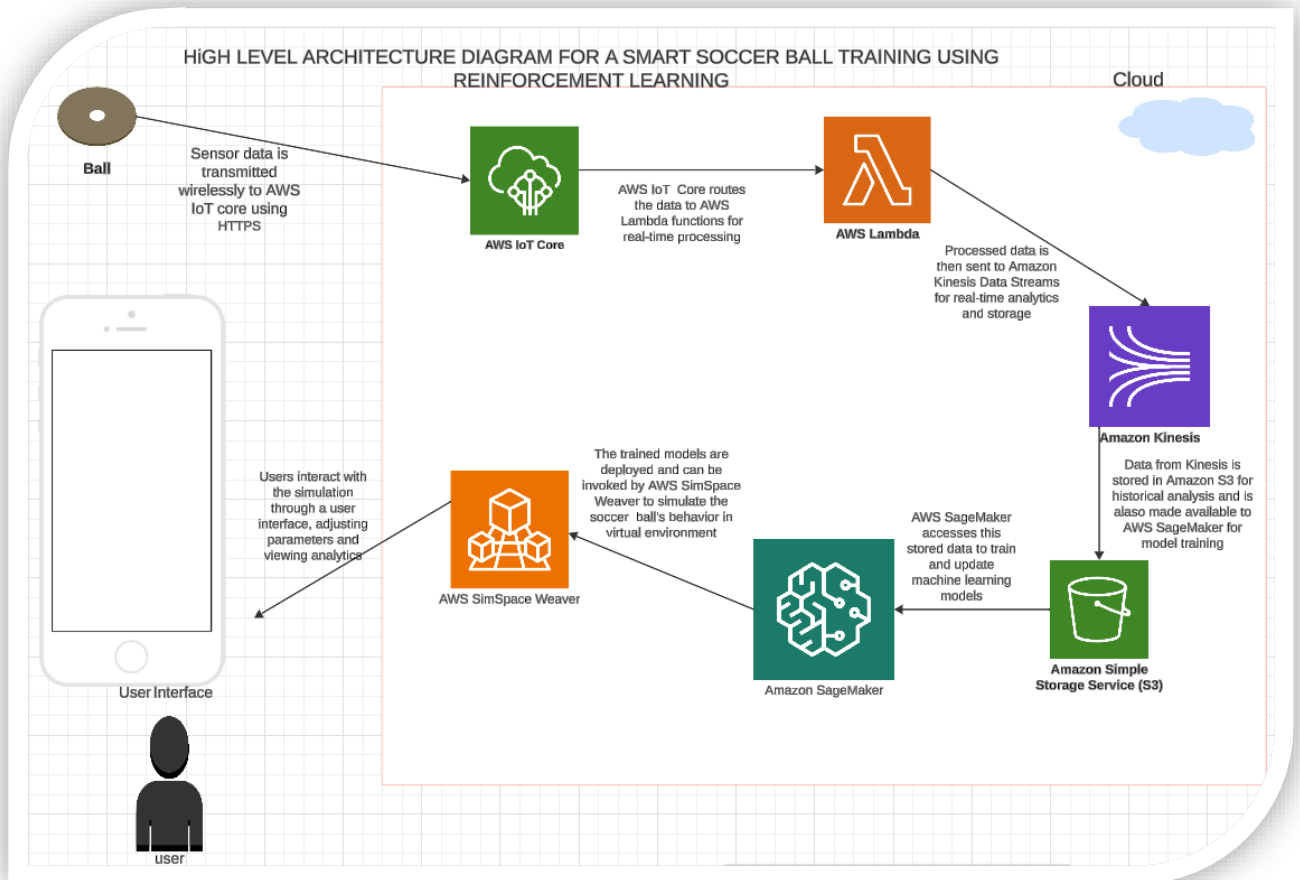
Zhang, Y. (2022). Genetic algorithms in sports: Enhancing soccer strategy through artificial intelligence. *Computers in Sport*

Science, 11(3), 215–231.
<https://doi.org/10.1016/j.coss.2022.07.003>

APPENDIX A

Architecture

This appendix presents the high-level architecture diagram outlining the integration of hardware, software, and cloud components in the Smart Soccer Ball project.



APPENDIX B

Local Environment Initialization, Agent Interaction, and Episode Completion Output

This appendix includes visualizations from TensorBoard logs, showcasing key training metrics such as reward progression and policy loss over time.

```
C:\Users\MissV\OneDrive\Documents\Education\CityU\2024FallQ4\AI620\Code> python environment.py
Custom observation shape: (1, 193) Observation space: Box(-inf, inf, (193,)), float32)
Action space: Discrete(7) C:\Users\MissV\OneDrive\Documents\Education\CityU\2024FallQ4\AI620\Code\venv\Lib\site-
packages\stable_baselines3\common\vec_env\base_vec_env.py:243: UserWarning: You tried to call render() but no render_mode
was passed to the env constructor. warnings.warn("You tried to call render() but no render_mode was passed to the env
constructor.")
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [0] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [0] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [0] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
```


[illegible]

[illegible]

APPENDIX C

Agent Training

This appendix provides a still image of the environment during training in human render mode, illustrating the agent's perspective within the grid-world maze.

rollout/		
ep_len_mean	12.9	
ep_rew_mean	0.955	
time/		
fps	723	
iterations	54	
time_elapsed	76	
total_timesteps	55296	
train/		
approx_kl	0.055399723	
clip_fraction	0.155	
clip_range	0.2	
entropy_loss	-0.167	
explained_variance	0.661	
learning_rate	0.0003	
loss	-0.0264	
n_updates	530	
policy_gradient_loss	-0.0259	
value_loss	0.000623	

rollout/		
ep_len_mean	12.2	
ep_rew_mean	0.957	
time/		
fps	723	
iterations	55	
time_elapsed	77	
total_timesteps	56320	
train/		
approx_kl	0.27952486	
clip_fraction	0.146	
clip_range	0.2	
entropy_loss	-0.129	
explained_variance	0.52	
learning_rate	0.0003	
loss	-0.0221	
n_updates	540	
policy_gradient_loss	-0.0178	
value_loss	0.000592	

rollout/		
ep_len_mean	11.9	
ep_rew_mean	0.958	
time/		
fps	720	
iterations	56	
time_elapsed	79	
total_timesteps	57344	
train/		
approx_kl	0.023819925	
clip_fraction	0.24	
clip_range	0.2	
entropy_loss	-0.217	
explained_variance	0.833	
learning_rate	0.0003	
loss	-0.0315	
n_updates	550	
policy_gradient_loss	-0.0296	
value_loss	0.000192	

rollout/		
ep_len_mean	12.3	
ep_rew_mean	0.957	
time/		

fps	720	
iterations	57	
time_elapsed	81	
total_timesteps	58368	
train/		
approx_kl	0.073045135	
clip_fraction	0.135	
clip_range	0.2	
entropy_loss	-0.154	
explained_variance	0.843	
learning_rate	0.0003	
loss	-0.04	
n_updates	560	
policy_gradient_loss	-0.0231	
value_loss	0.000152	

rollout/		
ep_len_mean	11.6	
ep_rew_mean	0.959	
time/		
fps	720	
iterations	58	
time_elapsed	82	
total_timesteps	59392	
train/		
approx_kl	0.06047496	
clip_fraction	0.109	
clip_range	0.2	
entropy_loss	-0.131	
explained_variance	0.589	
learning_rate	0.0003	
loss	0.00682	
n_updates	570	
policy_gradient_loss	-0.0187	
value_loss	0.000467	

Eval num_timesteps=60000, episode_reward=0.96 +/- 0.00
Episode length: 11.00 +/- 0.00

eval/		
mean_ep_length	11	
mean_reward	0.961	
time/		
total_timesteps	60000	
train/		
approx_kl	0.017942129	
clip_fraction	0.043	
clip_range	0.2	
entropy_loss	-0.094	
explained_variance	0.852	
learning_rate	0.0003	
loss	-0.0165	
n_updates	580	
policy_gradient_loss	-0.0123	
value_loss	0.000154	

rollout/		
ep_len_mean	13	
ep_rew_mean	0.954	
time/		
fps	718	
iterations	59	
time_elapsed	84	
total_timesteps	60416	

rollout/			explained_variance		
ep_len_mean			0.306		
ep_rew_mean			learning_rate		
time/			0.0003		
fps			loss		
iterations			-0.0343		
time_elapsed			n_updates		
total_timesteps			620		
train/			policy_gradient_loss		
approx_kl			-0.0198		
clip_fraction			value_loss		
clip_range			0.000805		
entropy_loss			-----		
explained_variance			rollout/		
learning_rate			ep_len_mean		
loss			15		
n_updates			ep_rew_mean		
policy_gradient_loss			0.947		
value_loss			time/		
-----			fps		
rollout/			719		
ep_len_mean			iterations		
ep_rew_mean			64		
time/			time_elapsed		
fps			91		
iterations			total_timesteps		
time_elapsed			65536		
total_timesteps			train/		
train/			approx_kl		
approx_kl			0.03838364		
clip_fraction			clip_fraction		
clip_range			0.217		
entropy_loss			clip_range		
explained_variance			0.2		
learning_rate			entropy_loss		
loss			-0.263		
n_updates			explained_variance		
policy_gradient_loss			0.667		
value_loss			learning_rate		
-----			0.0003		
rollout/			loss		
ep_len_mean			-0.073		
ep_rew_mean			n_updates		
time/			630		
fps			policy_gradient_loss		
iterations			-0.0305		
time_elapsed			value_loss		
total_timesteps			0.000518		
train/			-----		
approx_kl			rollout/		
clip_fraction			ep_len_mean		
clip_range			12.2		
entropy_loss			ep_rew_mean		
explained_variance			0.957		
learning_rate			time/		
loss			fps		
n_updates			718		
policy_gradient_loss			iterations		
value_loss			65		
-----			time_elapsed		
rollout/			92		
ep_len_mean			total_timesteps		
ep_rew_mean			66560		
time/			train/		
fps			approx_kl		
iterations			0.19488329		
time_elapsed			clip_fraction		
total_timesteps			0.347		
train/			clip_range		
approx_kl			0.2		
clip_fraction			entropy_loss		
clip_range			-0.282		
entropy_loss			explained_variance		
explained_variance			0.54		
learning_rate			learning_rate		
loss			0.0003		
n_updates			loss		
policy_gradient_loss			-0.0969		
value_loss			n_updates		
-----			640		
rollout/			policy_gradient_loss		
ep_len_mean			-0.0641		
ep_rew_mean			value_loss		
time/			0.00129		
fps			-----		
iterations			rollout/		
time_elapsed			ep_len_mean		
total_timesteps			12.6		
train/			ep_rew_mean		
approx_kl			0.956		
clip_fraction			time/		
clip_range			fps		
entropy_loss			718		
explained_variance			iterations		
learning_rate			66		
loss			time_elapsed		
n_updates			94		
policy_gradient_loss			total_timesteps		
value_loss			67584		
-----			train/		
rollout/			approx_kl		
ep_len_mean			0.12408055		
ep_rew_mean			clip_fraction		
time/			0.23		
fps			clip_range		
iterations			0.2		
time_elapsed			entropy_loss		
total_timesteps			-0.255		
train/			explained_variance		
approx_kl			0.888		
clip_fraction			learning_rate		
clip_range			0.0003		
entropy_loss			loss		
explained_variance			-0.0465		
learning_rate			n_updates		
loss			650		
n_updates			policy_gradient_loss		
policy_gradient_loss			0.0242		
value_loss			value_loss		
-----			4.94e-05		
rollout/			-----		
ep_len_mean			rollout/		
ep_rew_mean			ep_len_mean		
time/			11.7		
fps			ep_rew_mean		
iterations			0.959		
time_elapsed			time/		
total_timesteps			fps		
train/			717		
approx_kl					
clip_fraction					
clip_range					
entropy_loss					

	iterations		67	
	time_elapsed		95	
	total_timesteps		68608	
	train/			
	approx_kl		0.025122331	
	clip_fraction		0.173	
	clip_range		0.2	
	entropy_loss		-0.206	
	explained_variance		0.771	
	learning_rate		0.0003	
	loss		0.0226	
	n_updates		660	
	policy_gradient_loss		-0.0263	
	value_loss		0.000238	

	rollout/			
	ep_len_mean		11.9	
	ep_rew_mean		0.958	
	time/			
	fps		717	
	iterations		68	
	time_elapsed		97	
	total_timesteps		69632	
	train/			
	approx_kl		0.06556613	
	clip_fraction		0.141	
	clip_range		0.2	
	entropy_loss		-0.162	
	explained_variance		0.95	
	learning_rate		0.0003	
	loss		-0.0489	
	n_updates		670	
	policy_gradient_loss		-0.00874	
	value_loss		4.17e-05	

Eval num_timesteps=70000, episode_reward=0.96 +/- 0.00
 Episode length: 11.00 +/- 0.00

	eval/			
	mean_ep_length		11	
	mean_reward		0.961	
	time/			
	total_timesteps		70000	
	train/			
	approx_kl		0.062489584	
	clip_fraction		0.0564	
	clip_range		0.2	
	entropy_loss		-0.0847	
	explained_variance		0.778	
	learning_rate		0.0003	
	loss		-0.0217	
	n_updates		680	
	policy_gradient_loss		-0.0199	
	value_loss		0.000264	

	rollout/			
	ep_len_mean		11.1	
	ep_rew_mean		0.961	
	time/			
	fps		716	
	iterations		69	
	time_elapsed		98	
	total_timesteps		70656	

	rollout/			
	ep_len_mean		13.7	
	ep_rew_mean		0.952	
	time/			
	fps		715	
	iterations		70	
	time_elapsed		100	
	total_timesteps		71680	
	train/			

	approx_kl		0.05116283	
	clip_fraction		0.119	
	clip_range		0.2	
	entropy_loss		-0.131	
	explained_variance		0.985	
	learning_rate		0.0003	
	loss		-0.0654	
	n_updates		690	
	policy_gradient_loss		-0.0128	
	value_loss		9.36e-06	

	rollout/			
	ep_len_mean		13.2	
	ep_rew_mean		0.954	
	time/			
	fps		715	
	iterations		71	
	time_elapsed		101	
	total_timesteps		72704	
	train/			
	approx_kl		0.039104126	
	clip_fraction		0.158	
	clip_range		0.2	
	entropy_loss		-0.24	
	explained_variance		0.178	
	learning_rate		0.0003	
	loss		-0.0193	
	n_updates		700	
	policy_gradient_loss		-0.00839	
	value_loss		0.000591	

	rollout/			
	ep_len_mean		12.7	
	ep_rew_mean		0.956	
	time/			
	fps		715	
	iterations		72	
	time_elapsed		103	
	total_timesteps		73728	
	train/			
	approx_kl		0.058678307	
	clip_fraction		0.166	
	clip_range		0.2	
	entropy_loss		-0.236	
	explained_variance		0.761	
	learning_rate		0.0003	
	loss		0.0961	
	n_updates		710	
	policy_gradient_loss		-0.0254	
	value_loss		0.000299	

	rollout/			
	ep_len_mean		14.2	
	ep_rew_mean		0.95	
	time/			
	fps		714	
	iterations		73	
	time_elapsed		104	
	total_timesteps		74752	
	train/			
	approx_kl		0.042911883	
	clip_fraction		0.147	
	clip_range		0.2	
	entropy_loss		-0.215	
	explained_variance		0.893	
	learning_rate		0.0003	
	loss		-0.0327	
	n_updates		720	
	policy_gradient_loss		-0.0165	
	value_loss		0.000142	

	rollout/			
--	----------	--	--	--

	ep_len_mean		14.2	
	ep_rew_mean		0.95	
	time/			
	fps		710	
	iterations		74	
	time_elapsed		106	
	total_timesteps		75776	
	train/			
	approx_kl		0.02610638	
	clip_fraction		0.215	
	clip_range		0.2	
	entropy_loss		-0.269	
	explained_variance		0.556	
	learning_rate		0.0003	
	loss		-0.0248	
	n_updates		730	
	policy_gradient_loss		-0.0269	
	value_loss		0.000532	

	rollout/			
	ep_len_mean		13	
	ep_rew_mean		0.954	
	time/			
	fps		710	
	iterations		75	
	time_elapsed		108	
	total_timesteps		76800	
	train/			
	approx_kl		0.04112082	
	clip_fraction		0.249	
	clip_range		0.2	
	entropy_loss		-0.29	
	explained_variance		0.645	
	learning_rate		0.0003	
	loss		-0.0464	
	n_updates		740	
	policy_gradient_loss		-0.0367	
	value_loss		0.000662	

	rollout/			
	ep_len_mean		12.1	
	ep_rew_mean		0.958	
	time/			
	fps		711	
	iterations		76	
	time_elapsed		109	
	total_timesteps		77824	
	train/			
	approx_kl		0.058698382	
	clip_fraction		0.252	
	clip_range		0.2	
	entropy_loss		-0.229	
	explained_variance		0.883	
	learning_rate		0.0003	
	loss		-0.0224	
	n_updates		750	
	policy_gradient_loss		-0.0408	
	value_loss		0.000145	

	rollout/			
	ep_len_mean		12.5	
	ep_rew_mean		0.956	
	time/			
	fps		711	
	iterations		77	
	time_elapsed		110	
	total_timesteps		78848	
	train/			
	approx_kl		0.070883654	
	clip_fraction		0.226	
	clip_range		0.2	
	entropy_loss		-0.171	
	explained_variance		0.918	

	learning_rate		0.0003	
	loss		-0.0669	
	n_updates		760	
	policy_gradient_loss		-0.0423	
	value_loss		0.000216	

	rollout/			
	ep_len_mean		12.2	
	ep_rew_mean		0.957	
	time/			
	fps		712	
	iterations		78	
	time_elapsed		112	
	total_timesteps		79872	
	train/			
	approx_kl		0.017188694	
	clip_fraction		0.0901	
	clip_range		0.2	
	entropy_loss		-0.171	
	explained_variance		0.66	
	learning_rate		0.0003	
	loss		-0.0169	
	n_updates		770	
	policy_gradient_loss		-0.0119	
	value_loss		0.000467	

Eval num_timesteps=80000, episode_reward=0.96 +/- 0.00
Episode length: 11.00 +/- 0.00

	eval/			
	mean_ep_length		11	
	mean_reward		0.961	
	time/			
	total_timesteps		80000	
	train/			
	approx_kl		0.04313474	
	clip_fraction		0.107	
	clip_range		0.2	
	entropy_loss		-0.213	
	explained_variance		0.829	
	learning_rate		0.0003	
	loss		0.00571	
	n_updates		780	
	policy_gradient_loss		-0.0187	
	value_loss		0.000195	

	rollout/			
	ep_len_mean		12	
	ep_rew_mean		0.958	
	time/			
	fps		712	
	iterations		79	
	time_elapsed		113	
	total_timesteps		80896	

	rollout/			
	ep_len_mean		12.8	
	ep_rew_mean		0.955	
	time/			
	fps		711	
	iterations		80	
	time_elapsed		115	
	total_timesteps		81920	
	train/			
	approx_kl		0.06471266	
	clip_fraction		0.314	
	clip_range		0.2	
	entropy_loss		-0.294	
	explained_variance		0.938	
	learning_rate		0.0003	
	loss		-0.0486	
	n_updates		790	
	policy_gradient_loss		-0.0427	

value_loss	5.16e-05	

rollout/		
ep_len_mean	11.8	
ep_rew_mean	0.958	
time/		
fps	711	
iterations	81	
time_elapsed	116	
total_timesteps	82944	
train/		
approx_kl	0.06405732	
clip_fraction	0.211	
clip_range	0.2	
entropy_loss	-0.176	
explained_variance	0.795	
learning_rate	0.0003	
loss	-0.0243	
n_updates	800	
policy_gradient_loss	-0.0288	
value_loss	0.000221	

rollout/		
ep_len_mean	11.6	
ep_rew_mean	0.959	
time/		
fps	710	
iterations	82	
time_elapsed	118	
total_timesteps	83968	
train/		
approx_kl	0.061312027	
clip_fraction	0.0476	
clip_range	0.2	
entropy_loss	-0.0518	
explained_variance	0.881	
learning_rate	0.0003	
loss	-0.0393	
n_updates	810	
policy_gradient_loss	-0.023	
value_loss	0.000125	

rollout/		
ep_len_mean	12.2	
ep_rew_mean	0.957	
time/		
fps	708	
iterations	83	
time_elapsed	119	
total_timesteps	84992	
train/		
approx_kl	0.039377097	
clip_fraction	0.181	
clip_range	0.2	
entropy_loss	-0.117	
explained_variance	0.814	
learning_rate	0.0003	
loss	-0.0468	
n_updates	820	
policy_gradient_loss	0.022	
value_loss	0.000187	

rollout/		
ep_len_mean	11.8	
ep_rew_mean	0.958	
time/		
fps	708	
iterations	84	
time_elapsed	121	
total_timesteps	86016	
train/		
approx_kl	0.022505693	

clip_fraction	0.109	
clip_range	0.2	
entropy_loss	-0.158	
explained_variance	0.671	
learning_rate	0.0003	
loss	-0.0147	
n_updates	830	
policy_gradient_loss	-0.0127	
value_loss	0.000268	

rollout/		
ep_len_mean	15.4	
ep_rew_mean	0.946	
time/		
fps	707	
iterations	85	
time_elapsed	123	
total_timesteps	87040	
train/		
approx_kl	0.11171889	
clip_fraction	0.142	
clip_range	0.2	
entropy_loss	-0.131	
explained_variance	0.914	
learning_rate	0.0003	
loss	-0.0735	
n_updates	840	
policy_gradient_loss	-0.0356	
value_loss	6.9e-05	

rollout/		
ep_len_mean	14.4	
ep_rew_mean	0.95	
time/		
fps	706	
iterations	86	
time_elapsed	124	
total_timesteps	88064	
train/		
approx_kl	0.05634898	
clip_fraction	0.322	
clip_range	0.2	
entropy_loss	-0.321	
explained_variance	0.337	
learning_rate	0.0003	
loss	-0.0435	
n_updates	850	
policy_gradient_loss	-0.03	
value_loss	0.00163	

rollout/		
ep_len_mean	11.6	
ep_rew_mean	0.959	
time/		
fps	706	
iterations	87	
time_elapsed	126	
total_timesteps	89088	
train/		
approx_kl	0.1144023	
clip_fraction	0.191	
clip_range	0.2	
entropy_loss	-0.188	
explained_variance	0.632	
learning_rate	0.0003	
loss	-0.0365	
n_updates	860	
policy_gradient_loss	-0.0271	
value_loss	0.000624	

Eval num_timesteps=90000, episode_reward=0.96 +/- 0.00
Episode length: 11.00 +/- 0.00

eval/		
mean_ep_length	11	
mean_reward	0.961	
time/		
total_timesteps	90000	
train/		
approx_kl	0.065284915	
clip_fraction	0.063	
clip_range	0.2	
entropy_loss	-0.0858	
explained_variance	0.89	
learning_rate	0.0003	
loss	-0.0338	
n_updates	870	
policy_gradient_loss	0.0276	
value_loss	0.000171	

rollout/		
ep_len_mean	12.8	
ep_rew_mean	0.955	
time/		
fps	705	
iterations	88	
time_elapsed	127	
total_timesteps	90112	

rollout/		
ep_len_mean	14.6	
ep_rew_mean	0.949	
time/		
fps	704	
iterations	89	
time_elapsed	129	
total_timesteps	91136	
train/		
approx_kl	0.096123055	
clip_fraction	0.196	
clip_range	0.2	
entropy_loss	-0.191	
explained_variance	0.633	
learning_rate	0.0003	
loss	-0.0416	
n_updates	880	
policy_gradient_loss	-0.0249	
value_loss	0.000407	

rollout/		
ep_len_mean	11.1	
ep_rew_mean	0.961	
time/		
fps	704	
iterations	90	
time_elapsed	130	
total_timesteps	92160	
train/		
approx_kl	0.19162205	
clip_fraction	0.145	
clip_range	0.2	
entropy_loss	-0.15	
explained_variance	0.256	
learning_rate	0.0003	
loss	0.0508	
n_updates	890	
policy_gradient_loss	-0.0202	
value_loss	0.00348	

rollout/		
ep_len_mean	12.3	
ep_rew_mean	0.957	
time/		
fps	703	
iterations	91	

time_elapsed	132	
total_timesteps	93184	
train/		
approx_kl	0.06836694	
clip_fraction	0.347	
clip_range	0.2	
entropy_loss	-0.264	
explained_variance	0.581	
learning_rate	0.0003	
loss	-0.0699	
n_updates	900	
policy_gradient_loss	0.066	
value_loss	0.000154	

rollout/		
ep_len_mean	12.4	
ep_rew_mean	0.957	
time/		
fps	703	
iterations	92	
time_elapsed	133	
total_timesteps	94208	
train/		
approx_kl	0.020714343	
clip_fraction	0.161	
clip_range	0.2	
entropy_loss	-0.218	
explained_variance	0.581	
learning_rate	0.0003	
loss	-0.01	
n_updates	910	
policy_gradient_loss	-0.00672	
value_loss	0.000254	

rollout/		
ep_len_mean	12.3	
ep_rew_mean	0.957	
time/		
fps	703	
iterations	93	
time_elapsed	135	
total_timesteps	95232	
train/		
approx_kl	0.037652392	
clip_fraction	0.184	
clip_range	0.2	
entropy_loss	-0.199	
explained_variance	0.896	
learning_rate	0.0003	
loss	-0.0423	
n_updates	920	
policy_gradient_loss	-0.0175	
value_loss	0.000112	

rollout/		
ep_len_mean	11.2	
ep_rew_mean	0.961	
time/		
fps	702	
iterations	94	
time_elapsed	136	
total_timesteps	96256	
train/		
approx_kl	0.080989845	
clip_fraction	0.205	
clip_range	0.2	
entropy_loss	-0.133	
explained_variance	0.918	
learning_rate	0.0003	
loss	-0.05	
n_updates	930	
policy_gradient_loss	-0.0212	
value_loss	8.25e-05	

```

-----
| rollout/          |      |
| ep_len_mean      | 21    |
| ep_rew_mean      | 0.923 |
| time/            |      |
| fps              | 703   |
| iterations        | 95    |
| time_elapsed      | 138   |
| total_timesteps   | 97280 |
| train/           |      |
| approx_kl         | 0.13105541 |
| clip_fraction     | 0.316 |
| clip_range        | 0.2    |
| entropy_loss      | -0.254 |
| explained_variance | 0.971 |
| learning_rate      | 0.0003 |
| loss              | -0.0671 |
| n_updates         | 940    |
| policy_gradient_loss | 0.197 |
| value_loss        | 1.27e-05 |
-----

```

```

-----
| rollout/          |      |
| ep_len_mean      | 27.9  |
| ep_rew_mean      | 0.897 |
| time/            |      |
| fps              | 703   |
| iterations        | 96    |
| time_elapsed      | 139   |
| total_timesteps   | 98304 |
| train/           |      |
| approx_kl         | 0.1051268 |
| clip_fraction     | 0.194 |
| clip_range        | 0.2    |
| entropy_loss      | -0.144 |
| explained_variance | -0.119 |
| learning_rate      | 0.0003 |
| loss              | -0.0391 |
| n_updates         | 950    |
| policy_gradient_loss | 0.162 |
| value_loss        | 0.00478 |
-----

```

```

-----
| rollout/          |      |
| ep_len_mean      | 15.9  |
| ep_rew_mean      | 0.943 |
| time/            |      |
| fps              | 704   |
-----

```

```

-----
| iterations        | 97    |
| time_elapsed      | 140   |
| total_timesteps   | 99328 |
| train/           |      |
| approx_kl         | 0.5337625 |
| clip_fraction     | 0.366 |
| clip_range        | 0.2    |
| explained_variance | -0.0747 |
| learning_rate      | 0.0003 |
| learning_rate      | 0.0003 |
| loss              | -0.0265 |
| n_updates         | 960    |
| policy_gradient_loss | 0.308 |
| value_loss        | 0.00451 |
-----

```

Eval num_timesteps=100000, episode_reward=0.96 +/- 0.00
 Episode length: 11.00 +/- 0.00

```

-----
| eval/            |      |
| mean_ep_length    | 11    |
| mean_reward       | 0.961 |
| time/            |      |
| total_timesteps   | 100000 |
| train/           |      |
| approx_kl         | 0.13666381 |
| clip_fraction     | 0.168 |
| clip_range        | 0.2    |
| entropy_loss      | -0.155 |
| explained_variance | -0.427 |
| learning_rate      | 0.0003 |
| loss              | -0.0696 |
| n_updates         | 970    |
| policy_gradient_loss | 0.0638 |
| value_loss        | 0.00239 |
-----

```

```

-----
| rollout/          |      |
| ep_len_mean      | 13.1  |
| ep_rew_mean      | 0.954 |
| time/            |      |
| fps              | 705   |
| iterations        | 98    |
| time_elapsed      | 142   |
| total_timesteps   | 100352 |
-----

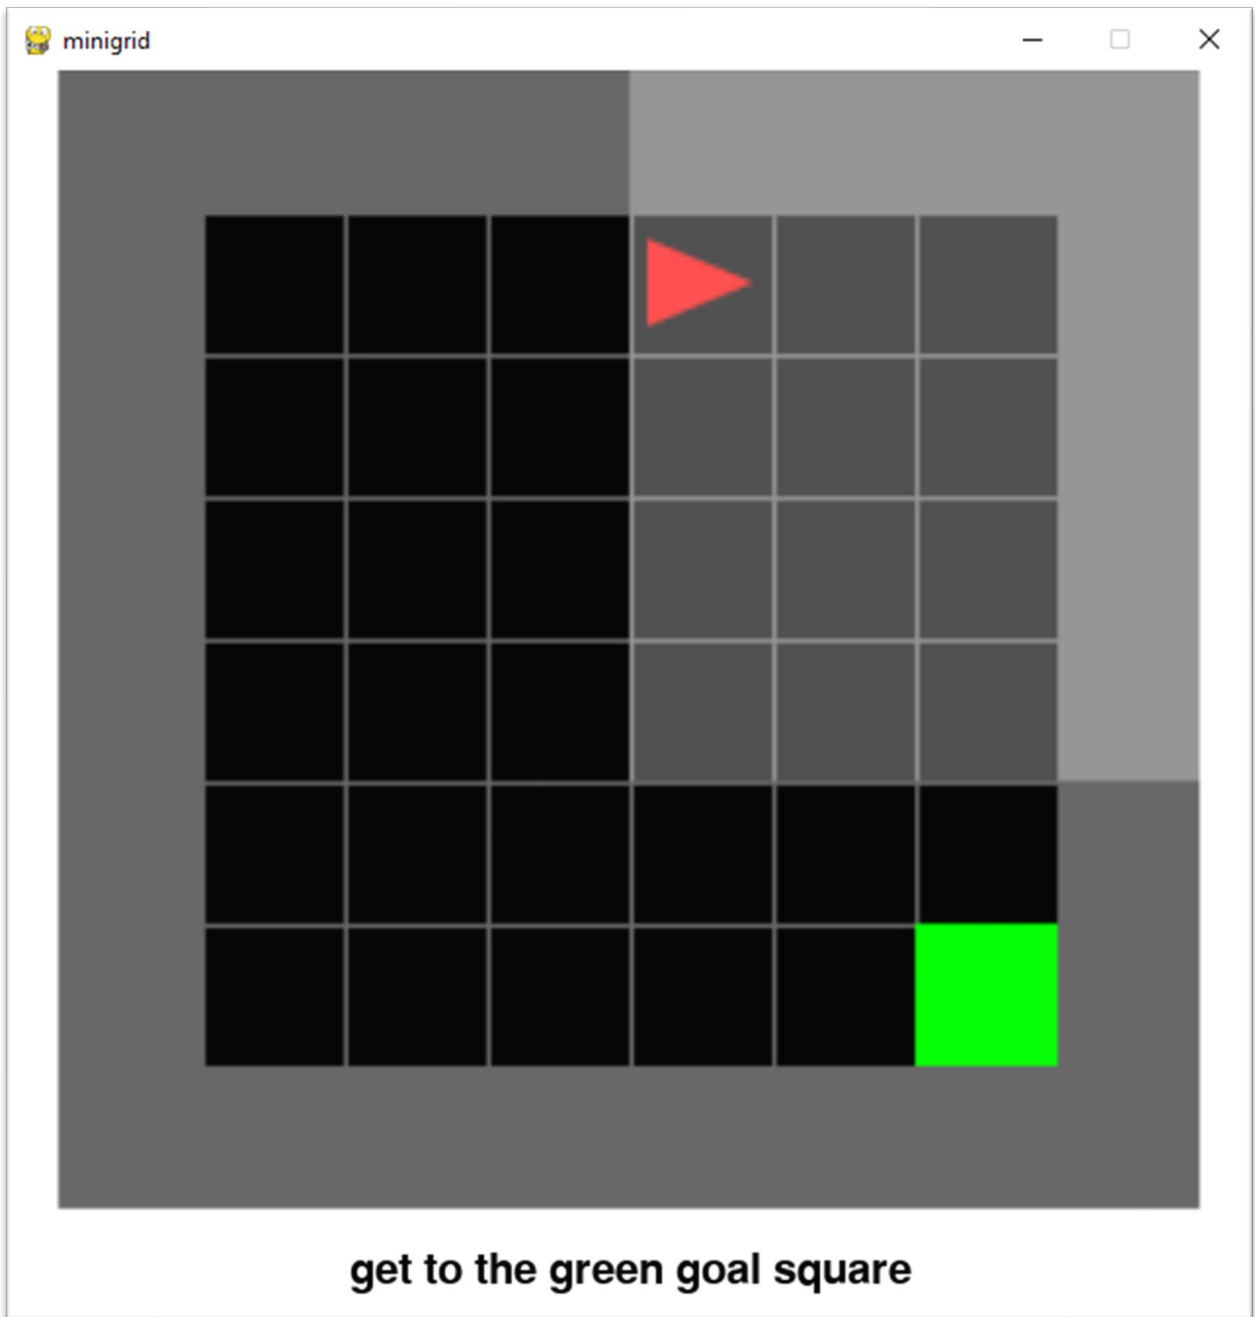
```

Training complete.
 Model saved as ppo_minigrid_model.
 Testing trained agent...

APPENDIX D

Render Mode = Human

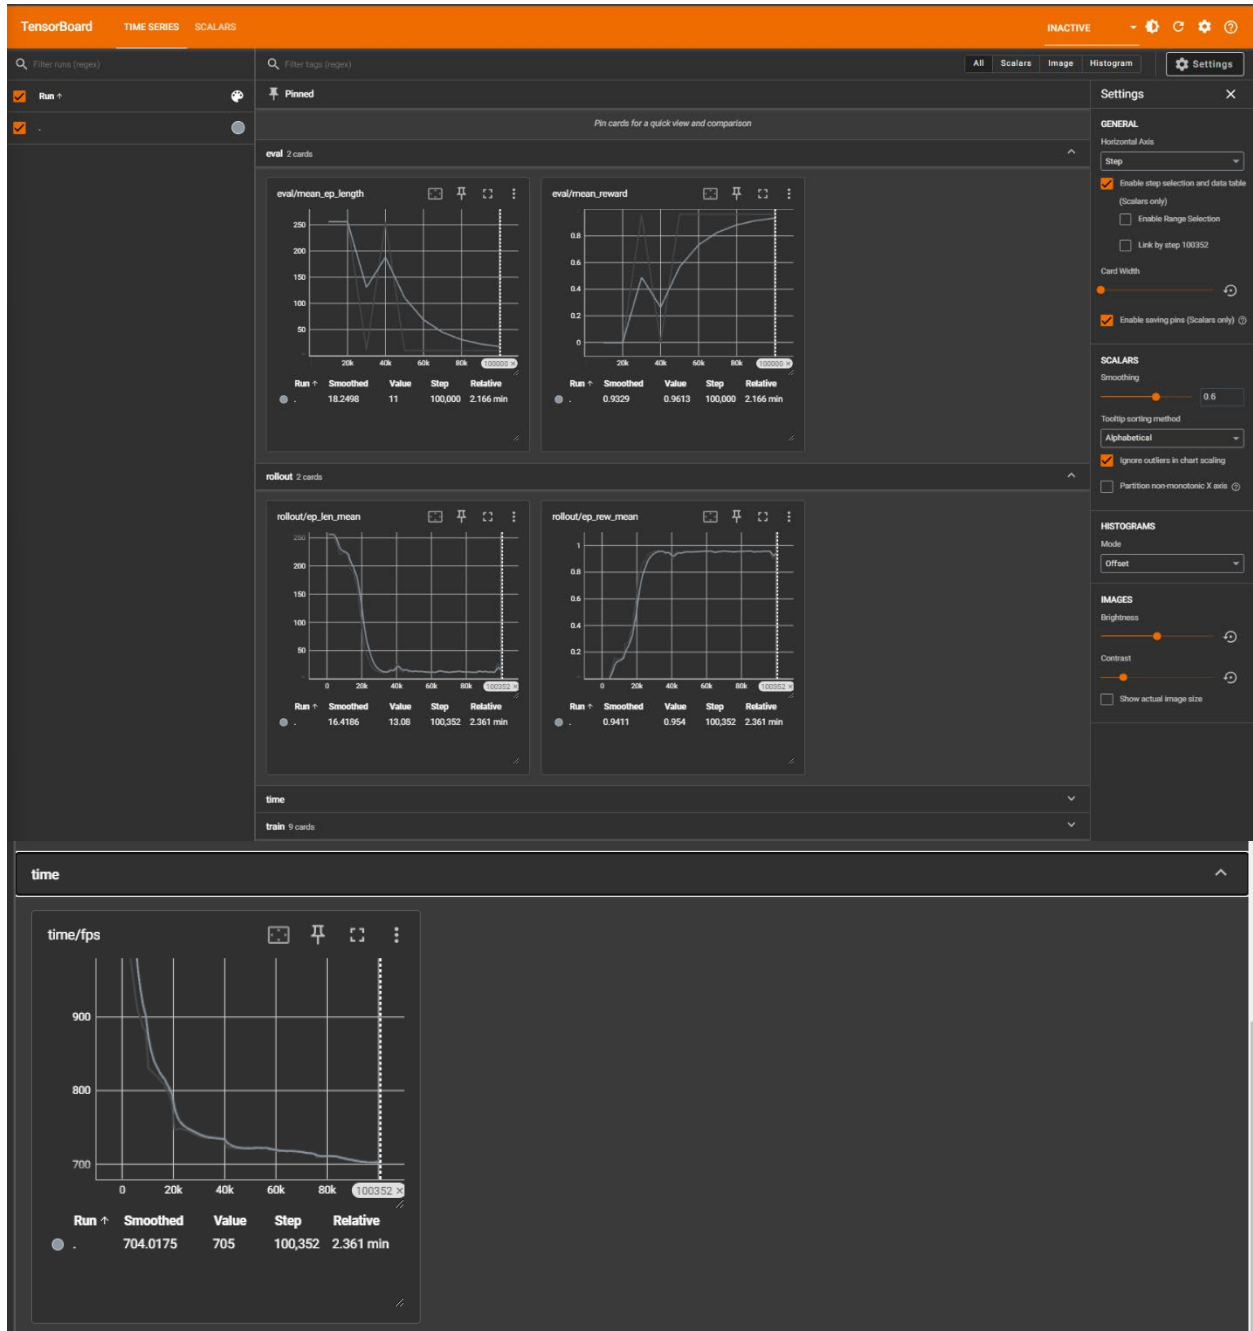
This still image, captured during training in human render mode, illustrates the agent (red triangle) navigating a grid environment toward the green goal square as per the specified task.

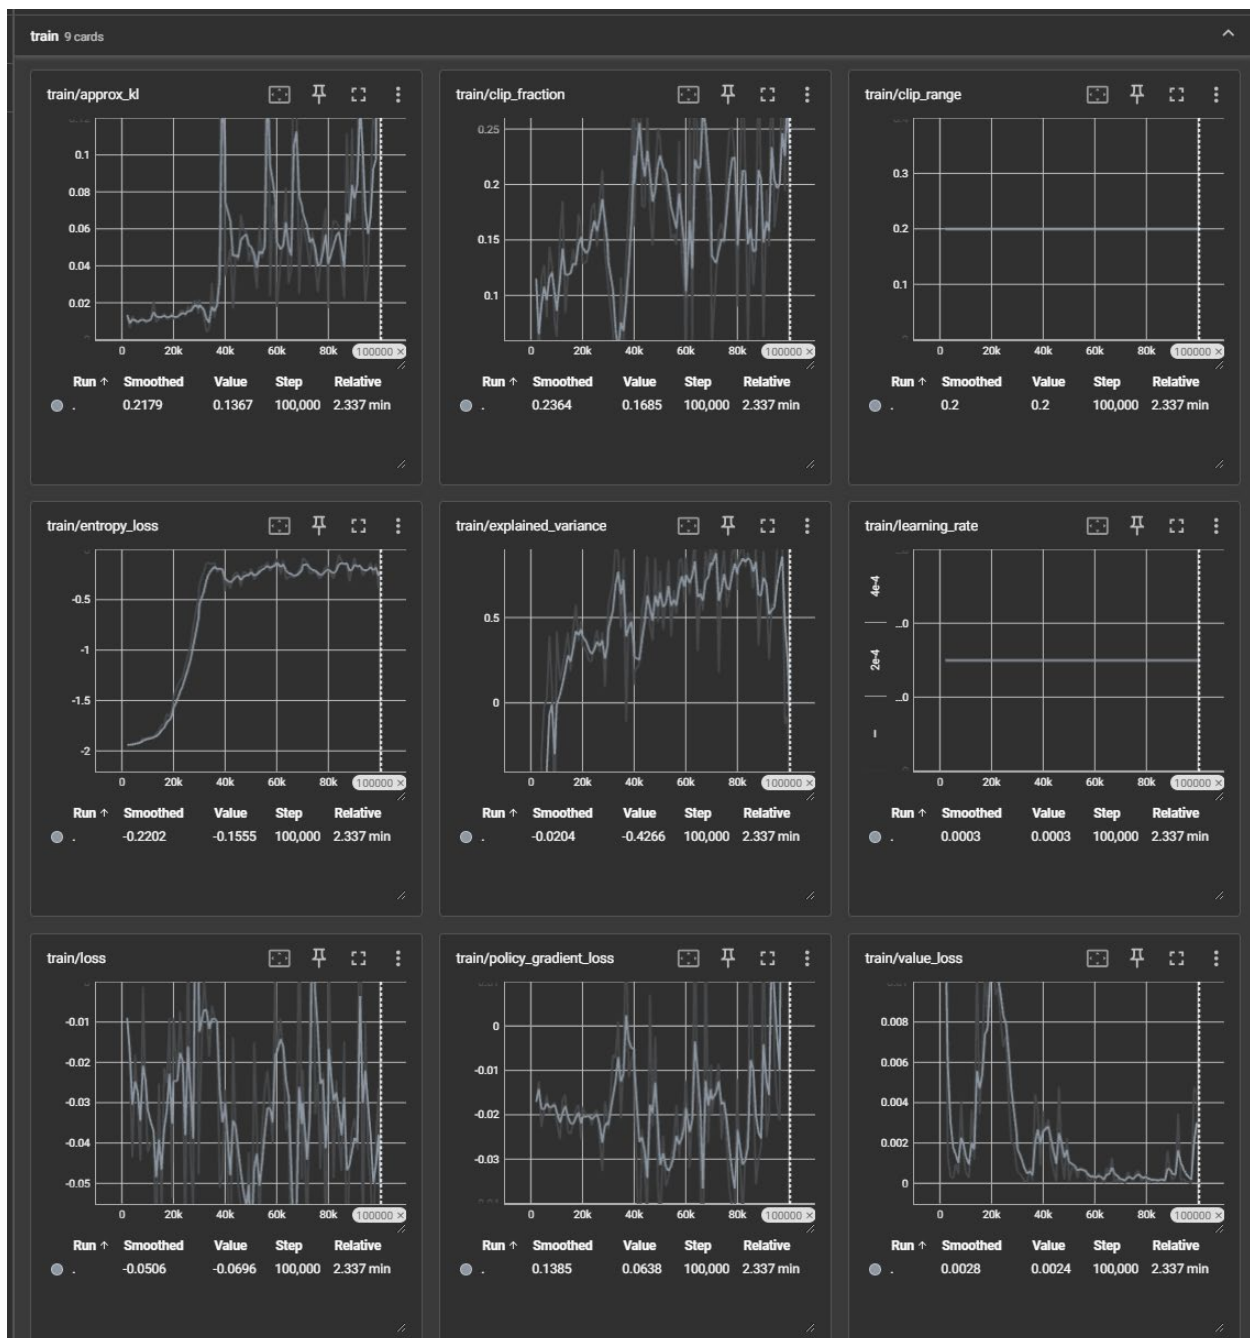


APPENDIX E

TensorBoard Logs – 1st Visualized Training Metrics

This appendix contains detailed logs from the initial evaluation phase, highlighting the agent's performance metrics, episode outcomes, and termination conditions during validation runs.





APPENDIX F

TensorBoard Logs – 2nd Visualized Training Metrics + Additional Details

This appendix provides comprehensive logs from the training runs, documenting the agent's progress, action sequences, rewards, and termination conditions to support the analysis of training performance.

